

Local Jarvis — Demo Walk-through

A 5-minute narrated demo you can run live or record. Run the guided script:

```
bash scripts/demo.sh
```

...or follow the steps manually below.

0. Health check (10 s)

```
jarvis doctor
```

Shows machine RAM, Ollama status, and which models are installed. Use this to confirm the environment is ready before demoing.

1. Local chat (30 s)

```
jarvis chat
```

Type a couple of questions to show fully-local, offline inference. `Ctrl-C` to exit.

2. Quantization snapshot — Part A (1 min)

For a quick demo, the committed `measurements.csv` already has results; just show the figures:

```
jarvis plots  
xdg-open benchmarks/plots/quant_size_vs_speed.png # or open on macOS
```

Explain the size ↔ speed ↔ quality trade-off and why **Q4_K_M** is the sweet spot.

Full run (slow): `jarvis bench-quant`

3. RAG vs no-RAG — Part C (1.5 min)

```
jarvis rag download # first time only  
jarvis rag index    # first time only  
jarvis rag compare
```

Point out how the RAG answer cites a source paper while the no-RAG answer is vaguer. Results saved to `outputs/rag_comparison.json`.

4. Web-search tool — Part D (1 min)

```
jarvis mcp-demo
```

Watch the model emit a `web_search` tool call, receive results, and synthesize an answer with a citation. Saved to `outputs/mcp_demo_results.json`.

5. Evaluation + report — Parts E & F (1 min)

```
jarvis eval      # 22-prompt suite, category-wise success rates
jarvis analyze   # builds outputs/analysis.md from all artifacts
```

Open `outputs/analysis.md` to show the data-driven limits section and the cloud-vs-local comparison.

Talking points

- **Privacy & cost:** everything but Part D runs offline; no per-token fees.
- **The knee of the curve:** Q4_K_M $\approx$ best quality-per-GB on 16 GB.
- **Honest limits:** show a failed adversarial/multi-step case from the eval.
- **Biggest lever:** a small GPU would 5–10 $\times$ throughput.